

Autonomous Job Application Agent

Problem Definition, Data Processing & Evaluation Criteria

1. Problem Definition

Tailoring a separate resume, cover letter, and interview-prep packet for every role is slow and repetitive. Off-the-shelf “AI resume writers” fail in two ways that matter for a real applicant: **(1) they fabricate** — inventing employers, titles, dates, metrics, or skills, which becomes a liability in an interview; and **(2) they leak personal data** — pasting a full resume into a third-party LLM exposes email, phone, address, and sometimes SSN / passport / credit-card numbers.

Problem statement

Build an autonomous, multi-agent system that, given a job description and a candidate’s resume, produces a tailored resume, cover letter, and interview-prep guide — without fabricating facts and without leaking sensitive personal data — while keeping a human in the loop for final approval.

Goals (priority order)

#	Goal	Why it matters
G1	No fabrication — every claim in generated docs traces back to the original resume	A false claim surfaces in interviews; trust is the core value
G2	PII protection — sensitive IDs never reach the LLM; standard PII redacted then restored in the final artifact	Privacy / regulatory exposure
G3	Relevance / JD alignment — reflect the role’s required skills, keywords, tone, seniority	The actual user value
G4	Human-in-the-loop — user reviews, gives per-document feedback, explicitly approves	User stays in control; the agent never auto-submits
G5	Quality — structured, concise, cliché-free, ATS-friendly documents	Determines competitiveness

Scope

In scope - Tailor an existing resume to one target JD (reframe, not rewrite) - Matched cover letter + interview-prep guide - Lightweight company research via public web search - PII detection, redaction, restoration - Human-feedback revision loop with per-document targeting - Export to Markdown / PDF	Out of scope (current version) - Auto-submitting to job boards / portals - Credential management or scraping behind logins - Authoring a resume from scratch - Multi-role batch tailoring in one run - Durable multi-user persistence (in-memory sessions)
---	---

Key constraints & assumptions

No fabrication: generation nodes may reframe / re-emphasize existing content but must not introduce new employers, titles, dates, metrics, or skills (prompt-enforced, verified in §3). **PII boundary:** analysis runs on a sanitized resume copy; raw PII is reattached only in the final artifact, and sensitive PII (SSN / credit-card / passport) aborts the run outright. **Assumptions:** the resume is truthful (the guarantee is relative to that input), and company research uses the free DuckDuckGo Instant Answer API (best-effort; a richer provider would swap in for production).

User interface

Users paste or upload the JD and resume, then review and approve the generated documents in the browser — a FastAPI web app (api.py, static/index.html) that streams node-by-node progress and collects per-document feedback.

2. Data Processing

Input	Req.	Formats	Handling
Job description	Yes	Text or PDF	PDF parsed to text; company name extracted by a cheap LLM call
Resume	Yes	Text / MD / PDF	Parsed via pypdf (layout-preserving); scanned / encrypted PDFs rejected; 10 MB cap

Processing pipeline (LangGraph state machine)

A 10-node LangGraph **StateGraph** (graph.py) over a shared **AgentState** TypedDict (state.py): parallel fan-out, conditional routing, an in-memory checkpointer, and an interrupt before human review.




Figure 1 — PII guard, research / analysis / writer sub-agents, the quality gate, and the human-review feedback loops (auto-retry, regenerate, re-ask).

PII data layer

Tier	Examples (regex-detected)	Action
Sensitive	SSN, credit-card, passport number	Abort the run (PIIBlockedError) — never sent to the LLM
Standard	email, phone, LinkedIn / GitHub URL, generic URL, address	Redact → token → restore

Flow: scan resume → replace each match with a reversible token [PII:TYPE:uid] → store pii_mapping in state → analysis & generation nodes see only the sanitized resume → restore tokens when assembling the final document. strip_stray_pii_tokens() removes any hallucinated token; logs are one-way redacted. The LLM never receives raw contact details or government identifiers.

State model (key fields)

Group	Fields
Inputs	job_description, resume_text, company_name
PII layer	pii_report (counts), pii_mapping, sanitized_resume
Research / analysis	company_research, jd_analysis
Generated docs	tailored_resume, cover_letter, interview_questions (+ sanitized copies for revision context)
Quality	quality_score, quality_feedback, revision_count
Feedback	human_feedback, feedback_type, approved, resume_feedback, cover_letter_feedback, interview_feedback

Feedback semantics & outputs

classify_feedback resolves input by priority: (1) the web-UI Approve flag is the only way to approve all three docs at once → END; (2) per-document feedback is relevance-checked against its target — irrelevant text is discarded, actionable text regenerates *only* that document. **revision_count** increments only on actionable feedback (gibberish never burns a cycle), and each node revises its own previous draft.

Outputs: tailored resume, cover letter, interview prep, and company research are written as Markdown under `outputs/{company}_*.md` and exportable to PDF via `/api/export-pdf`. Sessions are held in memory by `session_id` (a production build would move to Redis / SQLite).

Tech stack: LangGraph (≥ 0.2) · a configurable chat model via LangChain · FastAPI + Uvicorn + single-page UI · pypdf / fpdf2 · DuckDuckGo Instant Answer · deepeval + pytest.

3. Evaluation Criteria

A runnable suite under `evals/` drives the agent through its **real HTTP API** (in-process via Starlette TestClient, or a live server via `JOB_AGENT_BASE_URL`), so it exercises the exact handlers and background graph that production serves. It combines four layers: **LLM-as-judge** quality / grounding metrics, **deterministic guards** (no model, no billing), **regression controls** that pin the judges, and an **error-handling** suite for bad / unsafe inputs. The two golden examples (`backend_engineer_stripe`, `data_scientist_airbnb`) use resumes deliberately narrower than the JD, so the no-fabrication metrics have a real chance to catch invented skills or seniority.

3.1 Output quality & grounding (LLM-as-judge — deepeval)

Metric	Type	What it checks	Pass ≥
No Fabrication (resume)	GEval	Tailored resume invents no employer / title / date / metric / skill	0.80
Cover Letter No Fabrication	GEval	Every cover-letter claim about the candidate traces to the resume	0.80
Interview Prep No Fabrication	GEval	STAR answers grounded in the resume; generic Q&A not penalised	0.80
Faithfulness	deepeval	Tailored-resume claims grounded in the source resume (resume as context)	0.80
JD Alignment	GEval	Surfaces the JD-required skills / ATS keywords the candidate genuinely has	0.70
Cover Letter Quality	GEval	Specific hook, real experience tied to role, no clichés, concise	0.70

Thresholds follow goal priority (\$1): **no-fabrication & faithfulness gate at 0.80** (a violation is a hard failure); **relevance & quality gate at 0.70**. The judge tracks the agent’s model by default (`LLM_MODEL`); set `EVAL_MODEL` to pin an independent judge.

3.2 Deterministic guards (no model, no billing)

Guard	Checks	Budget
No stray PII tokens	Final docs contain no leftover [PII:~] redaction tokens (regex scan)	0 leaks
Per-node latency	No single LangGraph node exceeds the wall-clock budget (timed at fixture refresh)	≤ 60 s / node *
Token / cost budget	Total tokens per run within budget; optional \$ ceiling when a price is set	≤ 80k tok *

* Budgets are configurable via env (`EVAL_MAX_NODE_LATENCY_S`, `EVAL_MAX_TOKENS`, and `EVAL_MAX_COST_USD` / `EVAL_PRICE_PER_1M_*` for the \$ gate). Token counts are exact; the \$ gate is enforced only when a price is supplied, since the configured model’s rate isn’t known here.

3.3 No-fabrication regression controls (pin the judge itself)

`test_fabrication_regression.py` guards the detector against prompt / model drift using hand-built controls for *all three* documents (resume, cover letter, interview prep). For each, a blatantly **fabricated** version (invented Google staff role, team of 15, Stanford M.S., Go) must score **below** threshold, a **faithful** reframe must **pass**, and the fabricated version must score **strictly below** the faithful one — while generic technical Q&A is never flagged. No agent run is needed (the judge scores fixed text), but the judge still calls the model.

3.4 Error-handling / negative-path suite

`test_error_handling.py` grades behaviour on bad / unsafe inputs — does the agent fail *correctly*: right outcome, useful message, no data leak? Every case short-circuits **before any LLM call** (input validation at `/api/submit`, or the NODE 0 PII guard), so the whole file is deterministic, judge-free, and free to run on every push.

Case	Expected	Asserts
ssn_in_resume	status = error	Halts at the PII guard before any LLM; message says “Sensitive PII”; raw SSN never echoed; no output payload produced
empty_resume	HTTP 400	<code>/api/submit</code> rejects with “Resume text is required”
empty_jd	HTTP 400	<code>/api/submit</code> rejects with “Job description is required”
dashed SSN (guard unit)	detected	<code>scan_pii</code> flags a dashed SSN as sensitive

Success criteria

- ✓ No sensitive PII reached the LLM; final documents contain zero stray PII tokens.
- ✓ All no-fabrication / faithfulness metrics ≥ 0.80 ; JD-alignment & cover-letter quality ≥ 0.70 .
- ✓ Fabrication regression controls hold for resume, cover letter & interview prep (fabricated $<$ threshold $<$ faithful).
- ✓ Every adversarial input fails correctly (right outcome, useful message, no leak); per-node latency & token budgets met.
- ✓ Runtime quality gate reached ≥ 0.7 (or hit the revision cap) and the human explicitly approved the artifacts.

Known limitations & future work

Search: DuckDuckGo Instant Answer is shallow (swap for Tavily / SerpAPI). **Persistence:** in-memory sessions are lost on restart (move to Redis / SQLite). **PII detection** is regex-based — an NER detector would raise recall. **Dataset:** two golden examples; broader role coverage would strengthen the signal. **No auto-submission** by design — the agent stops at approved artifacts.